

Implementing GANs

Duke MLSS

Alex Lew

Probability Distributions

$$x \sim p(x)$$

x “is drawn from” $p(x)$

x “is a sample from” $p(x)$

- 
- $p(x)$ a distribution assigning 50% probability to H, 50% to T
 - x a random variable, equal to H or T

Probability Distributions

$$\textcolor{green}{x} \sim p(x)$$

$\textcolor{green}{x}$ “is drawn from” $p(x)$

$\textcolor{green}{x}$ “is a sample from” $p(x)$



- $p(\mathbf{x})$ a distribution assigning some small probability to each possible setting of pixel values in a 24-bit-color 100x100 image.
- \mathbf{x} a random variable, equal to some 100x100 image.

Generative modeling

Approximate $p(x)$ that generated the data

Do we want to...

- evaluate relative likelihoods of different data points?
- compute conditional distributions?
- do inference on latent variables?
- produce new samples (perhaps conditioned on something)?

Note: we have no “labeled” $p(x)$ values,
just lots of x

Generative modeling

For today:

Generate New Samples

(create “realistic” images, sounds, etc.)

Sampling

How do we generate a sample from a distribution?

Method 1: Go out and flip a coin.
(slow, expensive)

 **p(x)** a distribution assigning 50% probability to H, 50% to T
x a random variable, equal to H or T

Sampling

How do we generate a sample from a distribution?

Method 2: *Simulate* a sample.

```
np.random.uniform(0,1)
```



Source of (pseudo-)randomness

```
def G(z):  
    if z > 0.5:  
        return "Heads"  
    else:  
        return "Tails"
```



Deterministic mapping



Sample from
 $p(x)$

Sampling

But sometimes, we don't know exactly what $G(z)$ should be.



Sampling

But sometimes, we don't know exactly what $G(z)$ should be.
We can estimate *parameters* from *data*.

```
np.random.uniform(0,1)
```



Source of (pseudo-)randomness

```
def G(z):  
    if z > theta:  
        return "Heads"  
    else:  
        return "Tails"
```



Deterministic mapping



Sample from
 $p(x)$

Sampling

How do we generate a sample from a distribution?

Method 1: Go out and take a picture.
(slow, expensive)



- $p(x)$ a distribution assigning some small probability to each possible setting of pixel values in a 24-bit-color 100×100 image.
- x a random variable, equal to some 100×100 image.

Sampling

How do we generate a sample from a distribution?

Method 2: *Simulate* a sample.

```
np.random.uniform(0,1)
```



Source of (pseudo-)randomness

```
def G(z):  
    # hmmmm...  
    # help?
```



Deterministic mapping



Sample from
 $p(x)$

How to write \mathbf{G} ?

It's just a
function: learn it!

```
np.random.uniform(0,1)
```

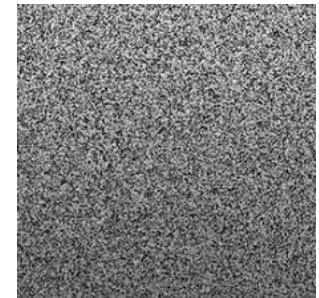


Source of (pseudo-)randomness

```
def G(z):  
    return MLP(z, theta)
```



Deterministic mapping with
randomly initialized parameters



Sample from
 $p(x)$

Learn \mathbf{G}

But with what loss function?

```
np.random.uniform(0,1)
```

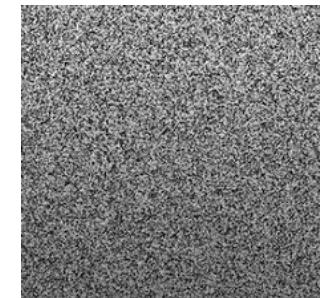


Source of (pseudo-)randomness

```
def G(z):  
    return MLP(z, theta)
```


$$G_{\theta}$$

Deterministic mapping with
randomly initialized parameters

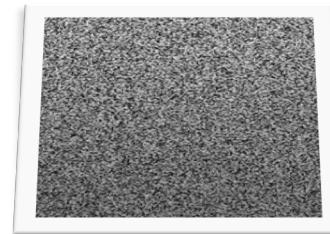


Sample from
 $p(x)$

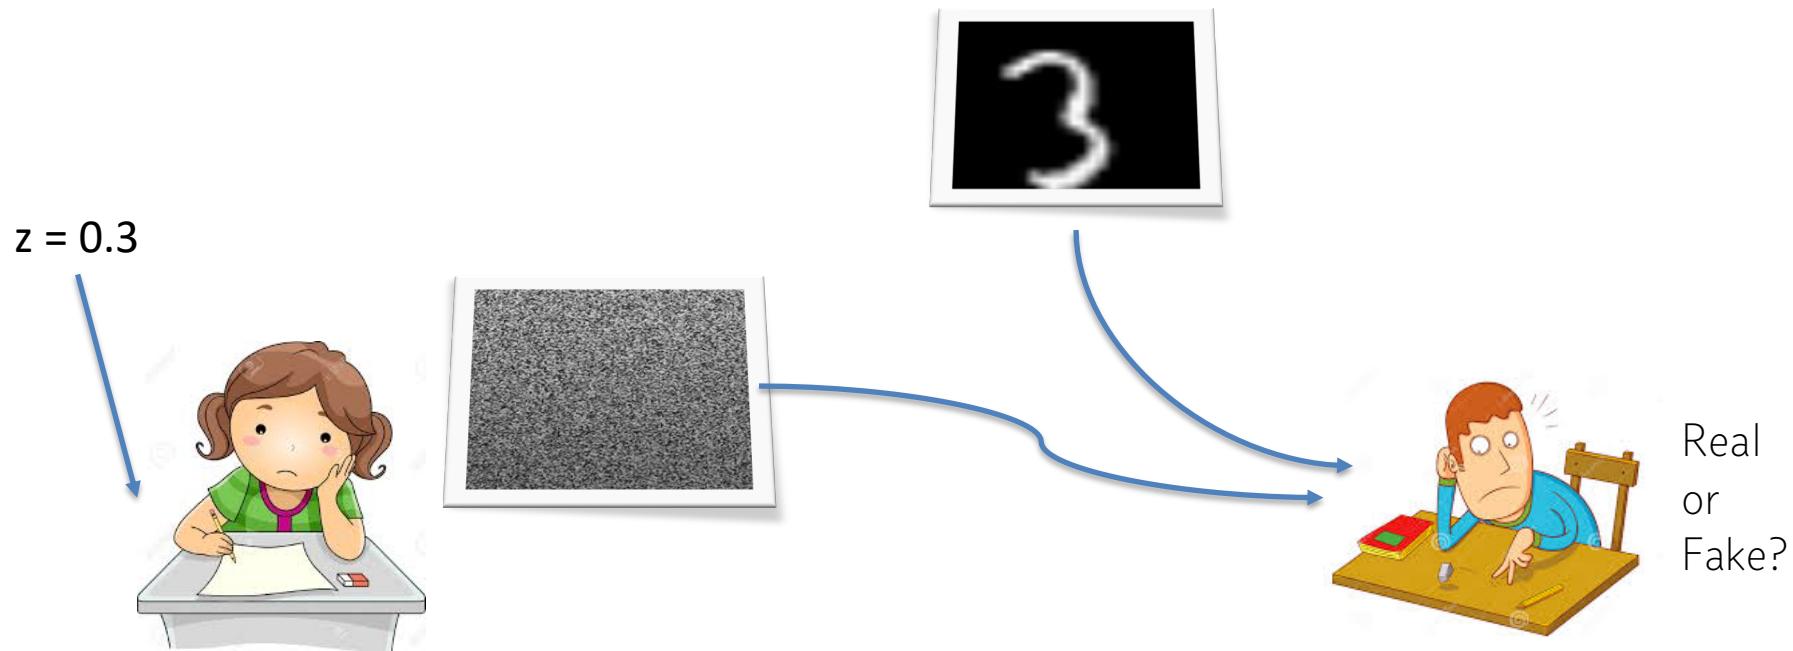
Learn **G**

How do you teach **G** with no (in, out) examples?

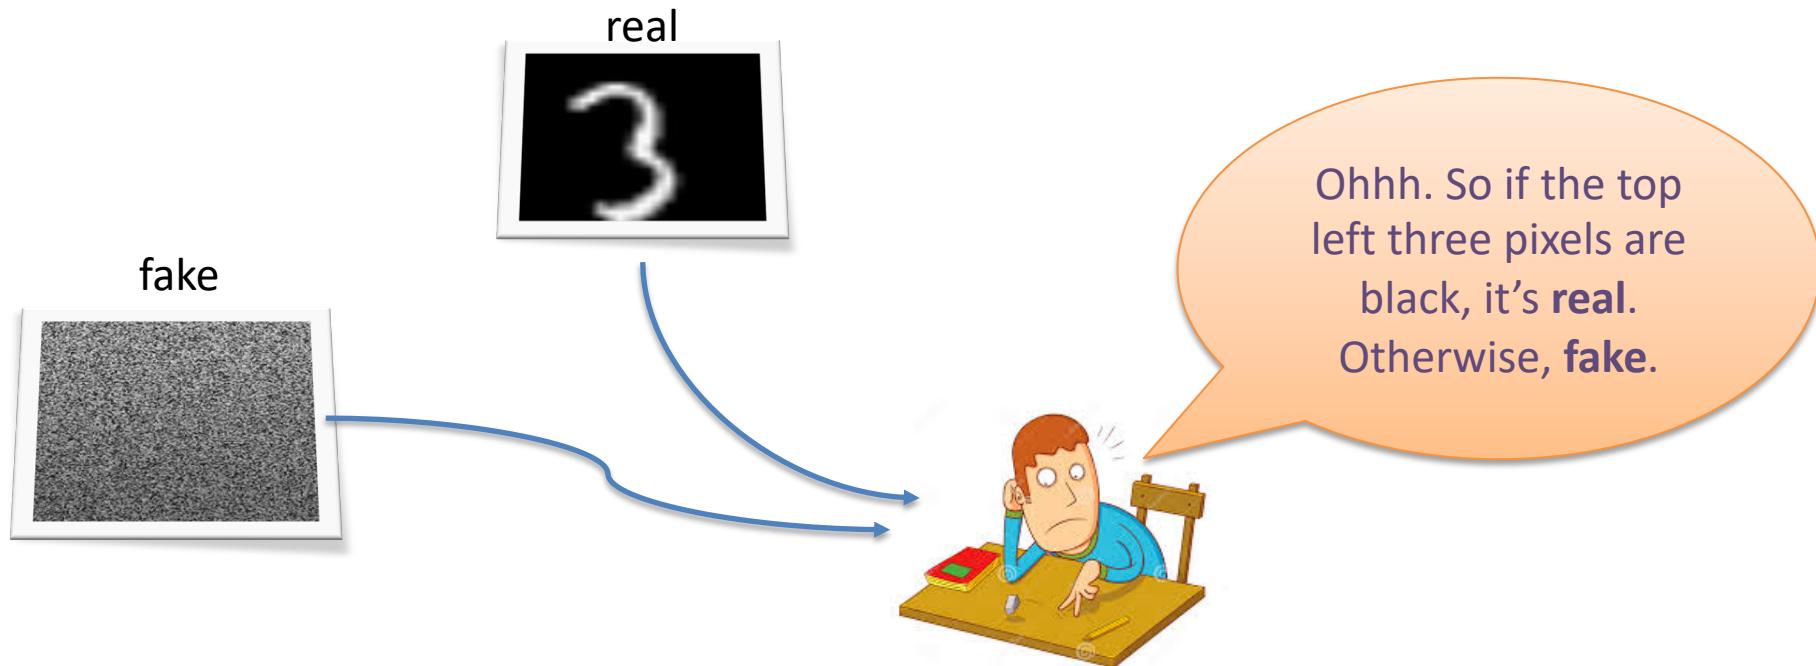
I know my answer isn't right, but what should it have been for $z=0.3$?



Get two students (neural networks)



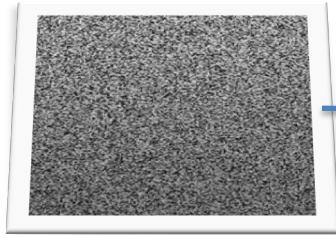
Get two students (neural networks)



Get two students (neural networks)

You're failing to
fool the
discriminator. Try
making the top left
pixels darker.

$z = 0.3$



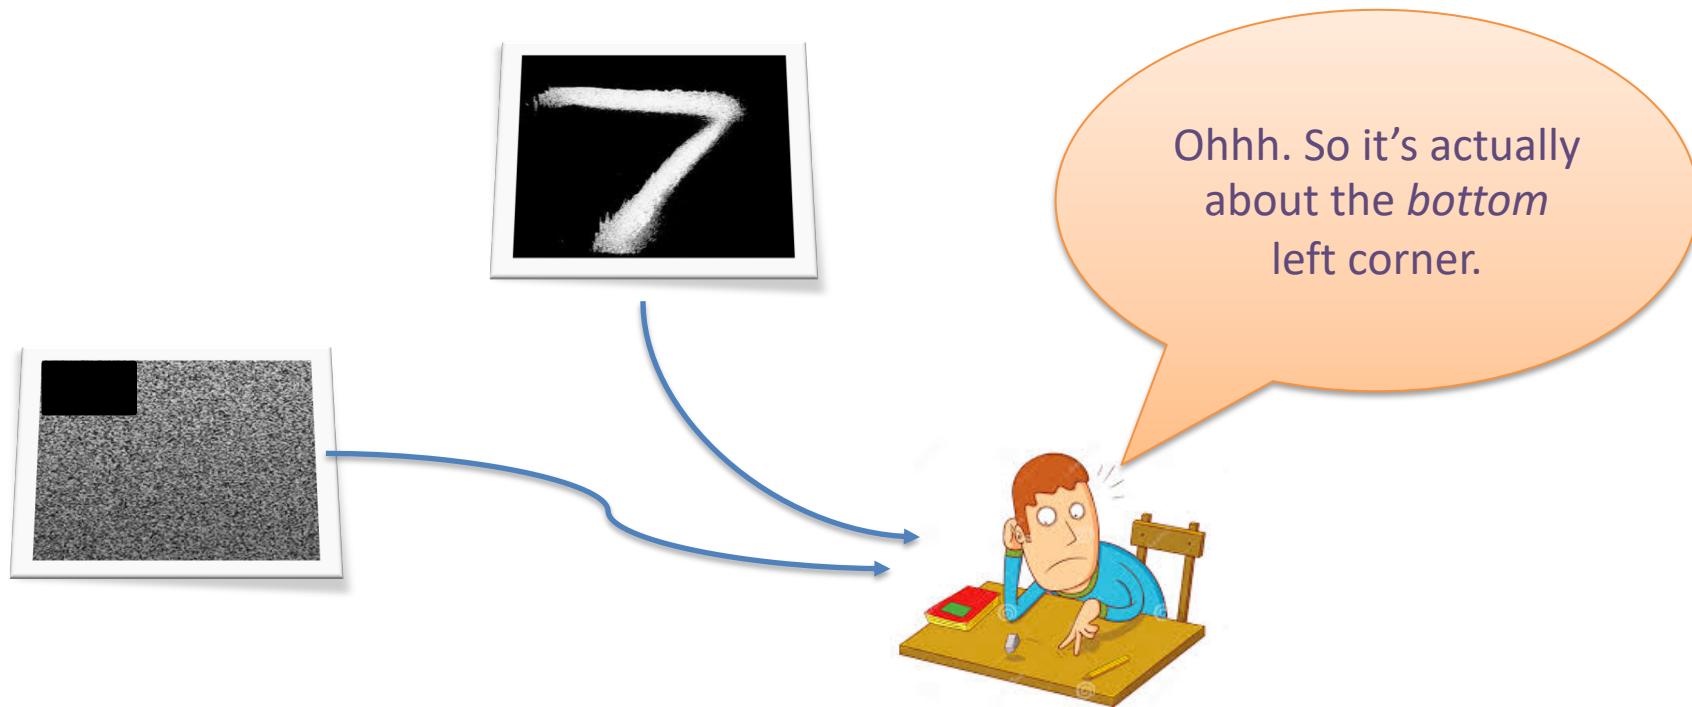
Get two students (neural networks)

You're failing to
fool the
discriminator. Try
making the top left
pixels darker.

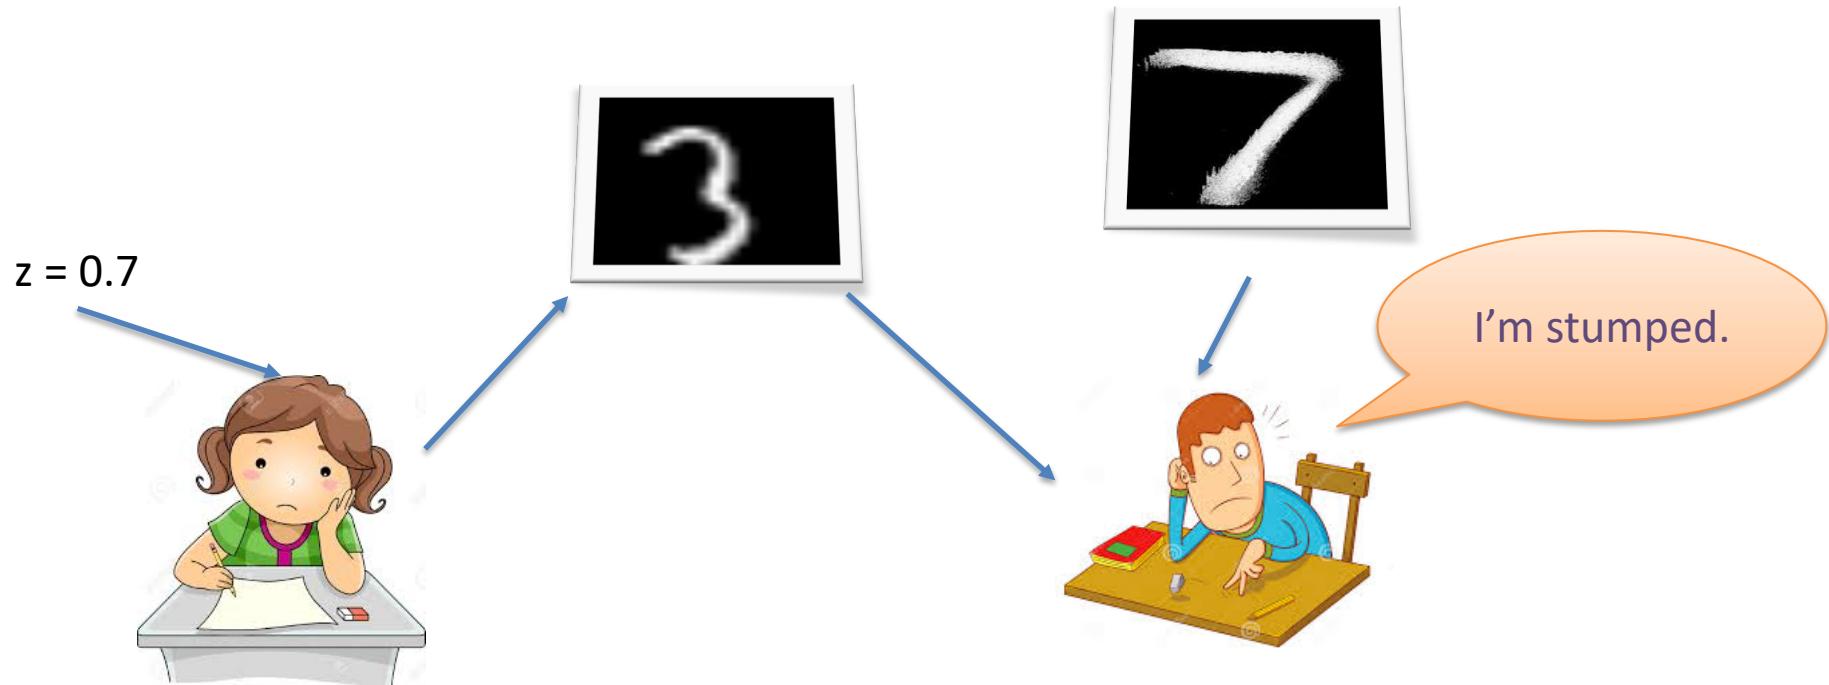
$z = 0.5$



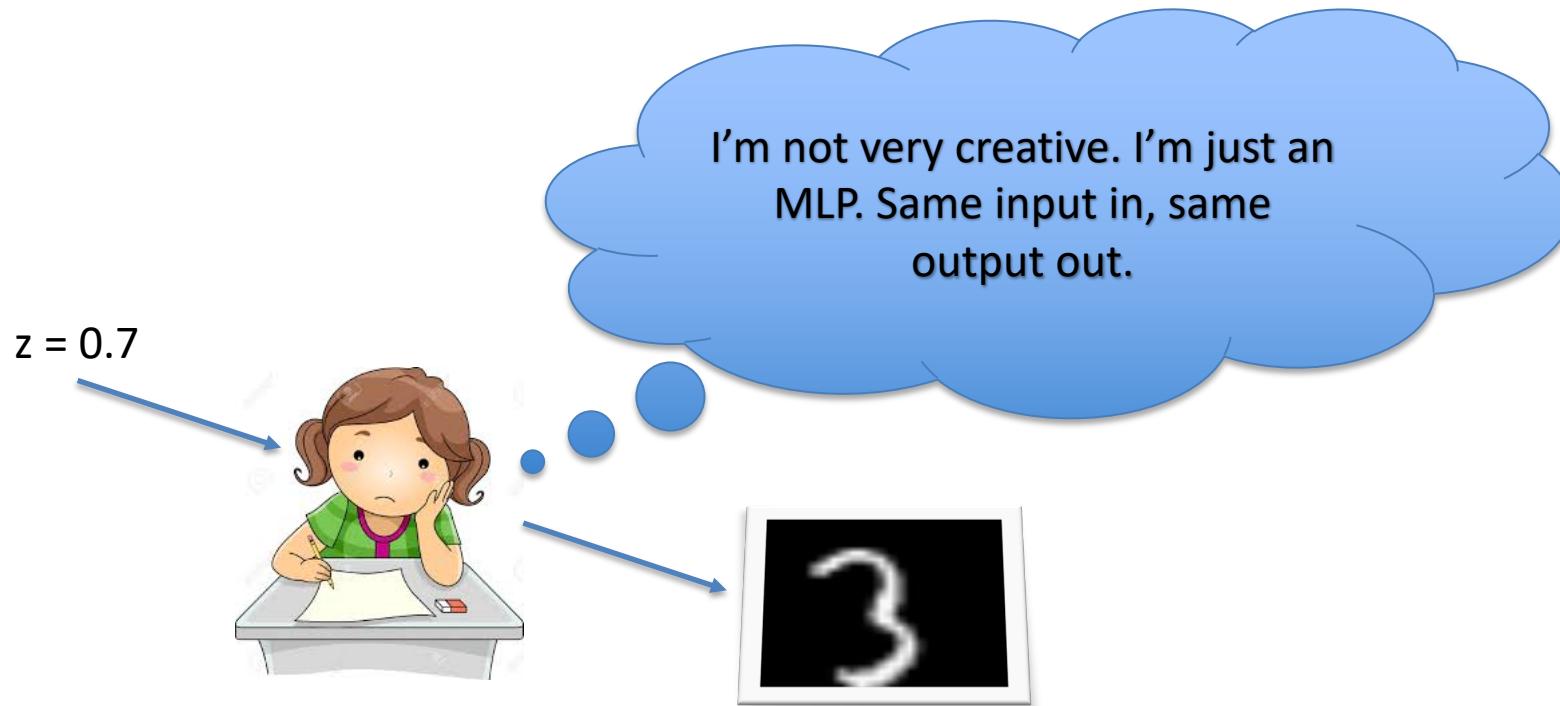
Get two students (neural networks)



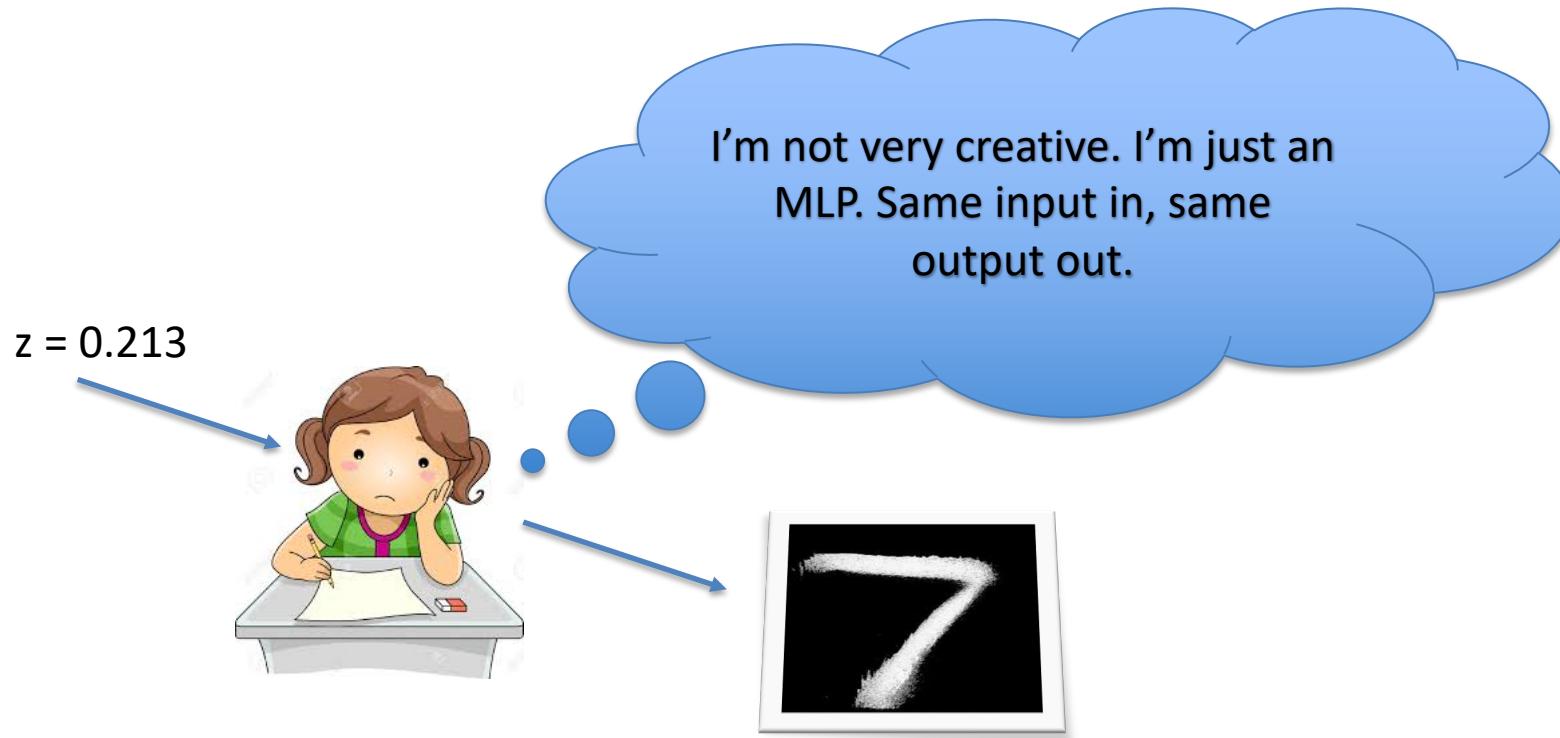
Many iterations later...



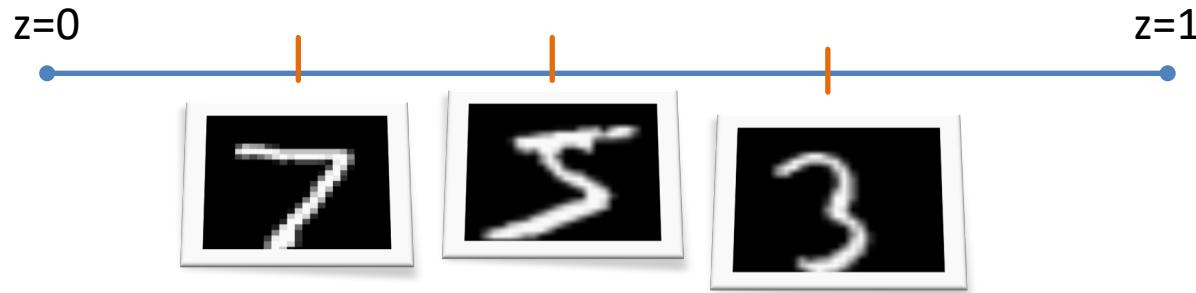
The importance of z



The importance of z

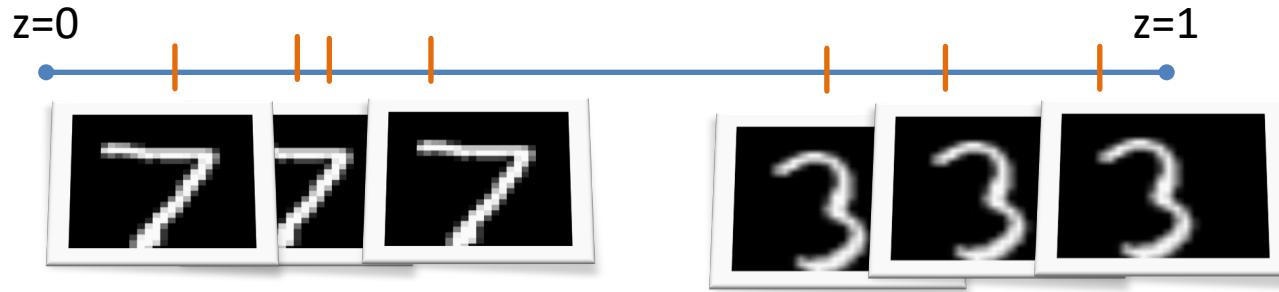


The importance of z



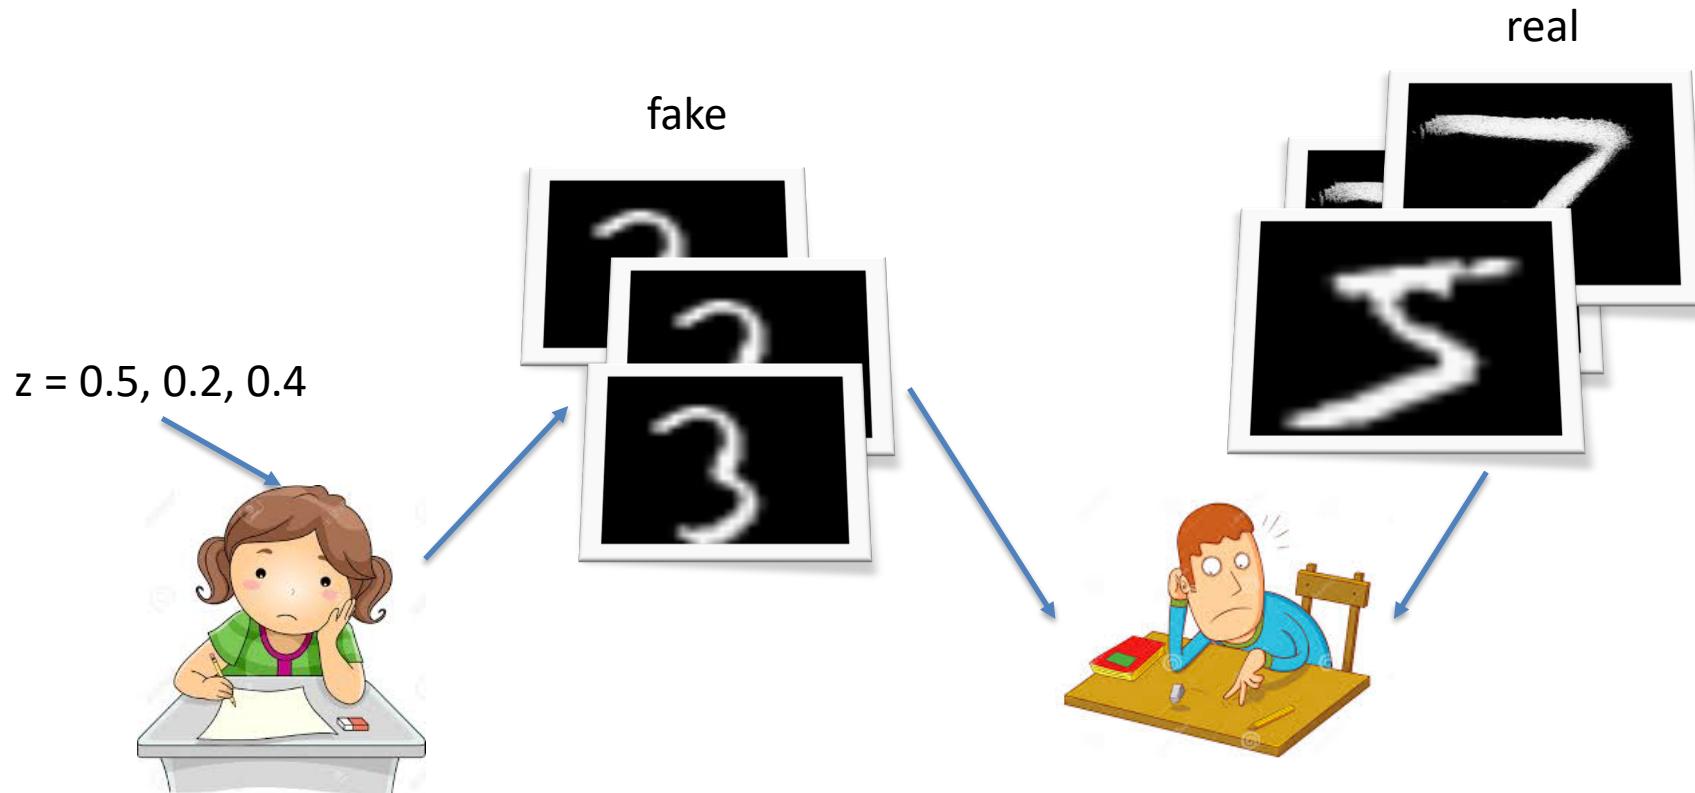
G a *deterministic function*
from z to images

The importance of z

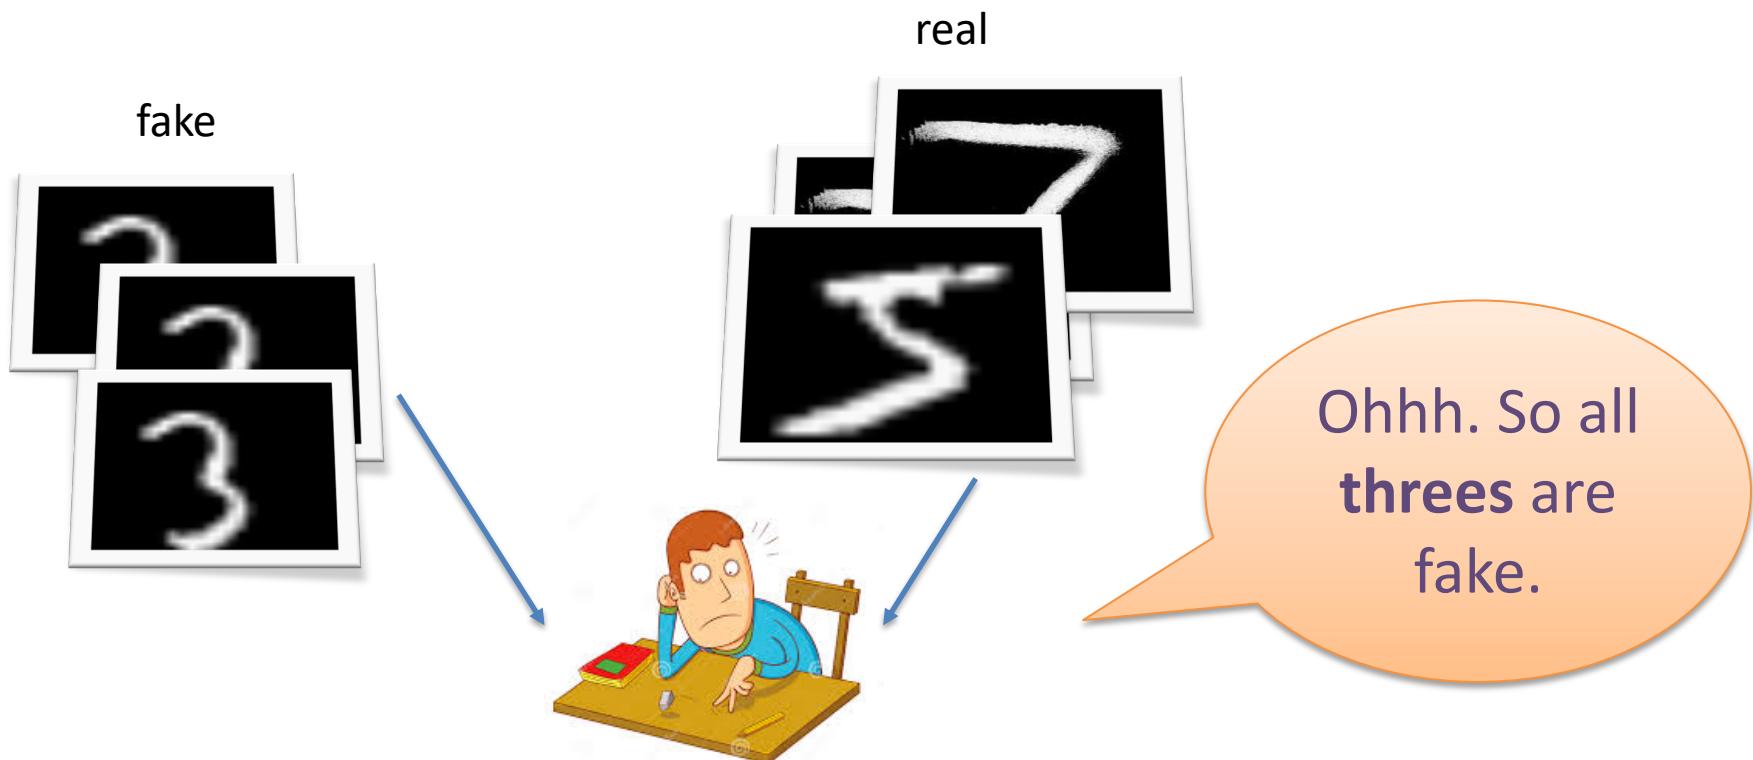


What if G learns to “play it safe” and always generate similar images?

What if G always writes the same thing?



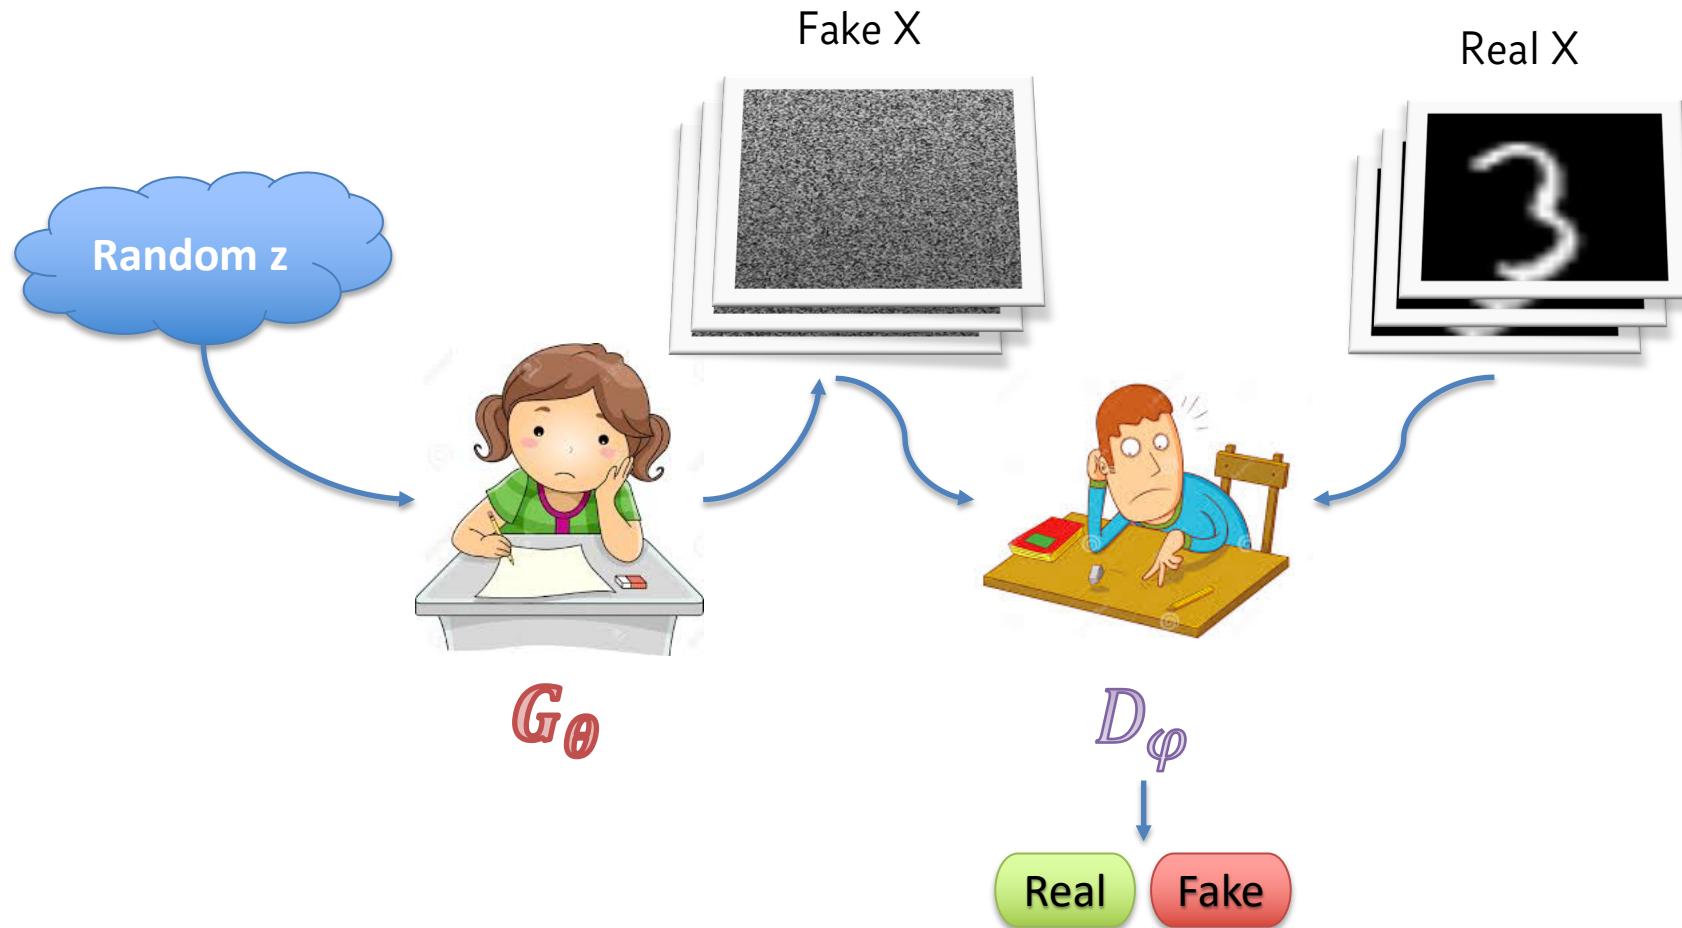
What if G always writes the same thing?



Main takeaway



We don't know how to teach G to draw. But with two students, we can make them teach each other!



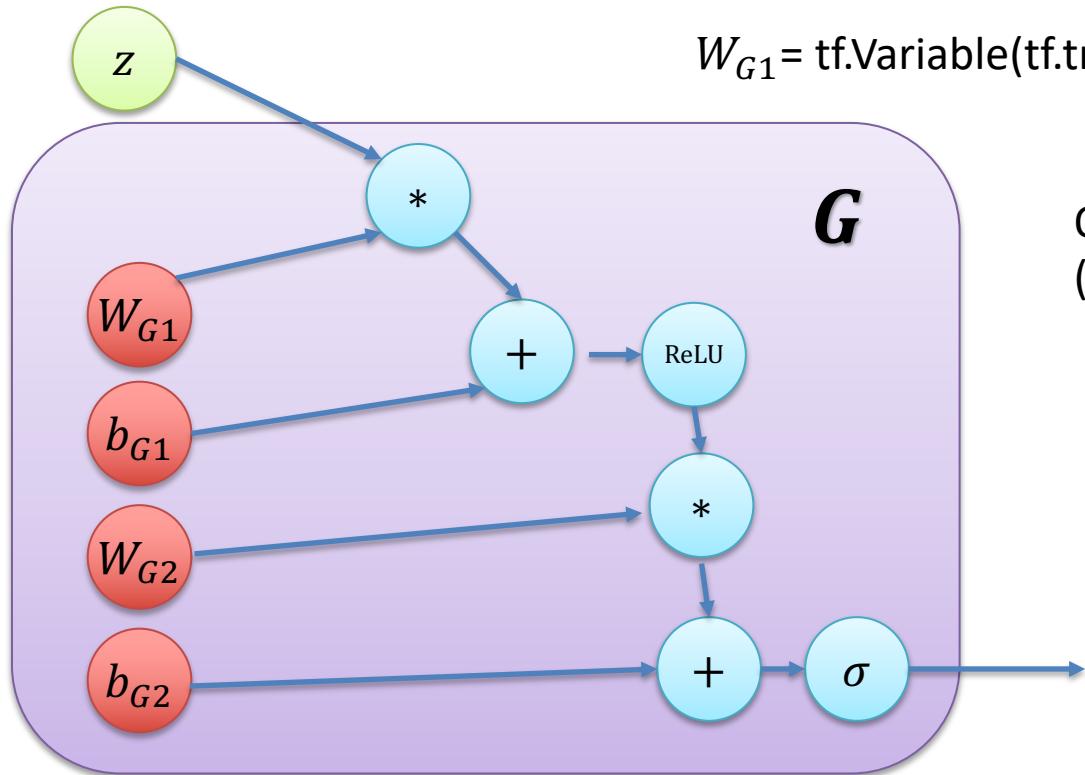
Implementation subtleties



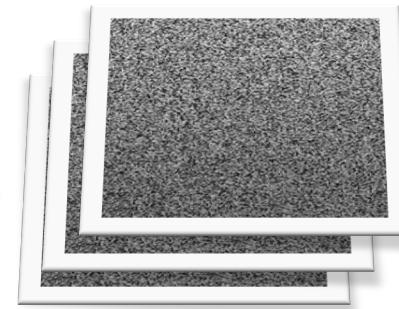
$z = \text{tf.placeholder}([50, 20])$

- We provide *multiple* random numbers (say, 20) to inform G in creating *each* image.
- We ask G to create a whole “batch” (say, 50) of fake images each batch.

Architecture of G : $20 \rightarrow 256 \rightarrow 784$



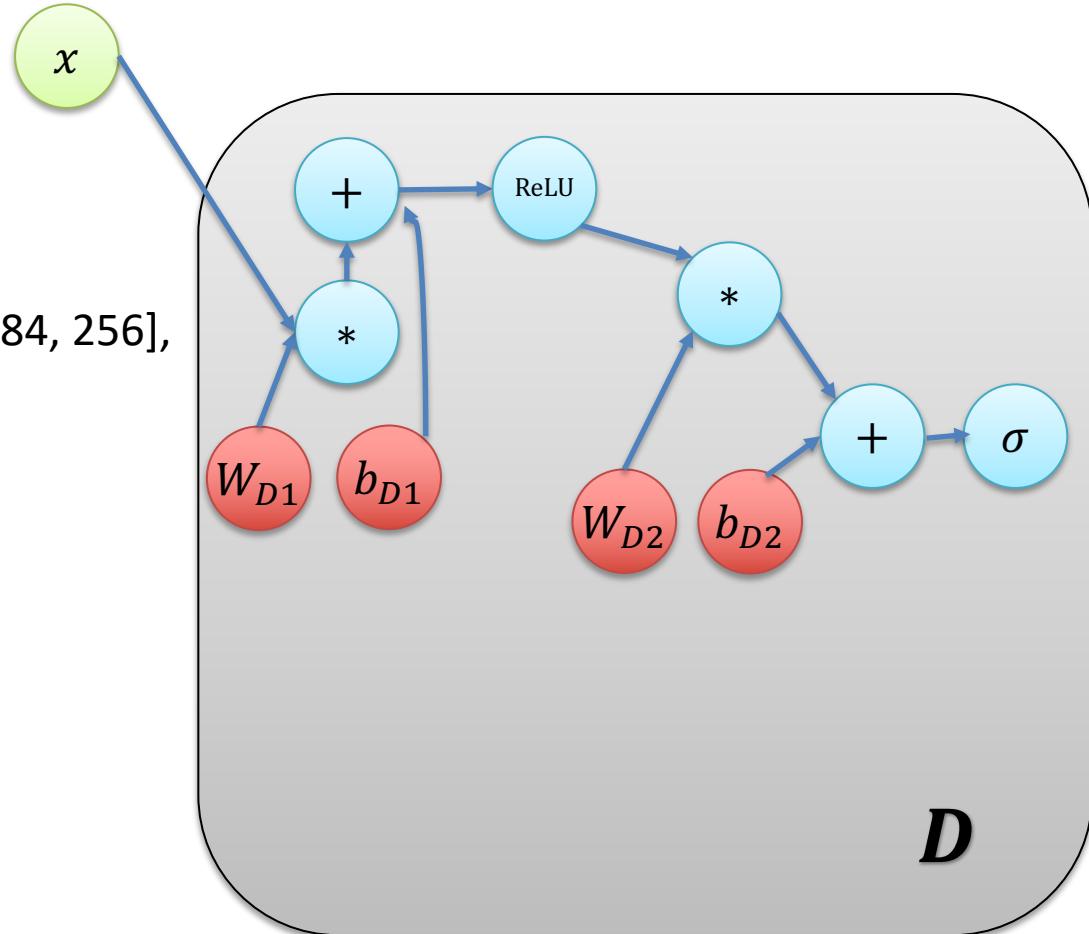
Our output: values between 0 and 1
(black and white) for each pixel.



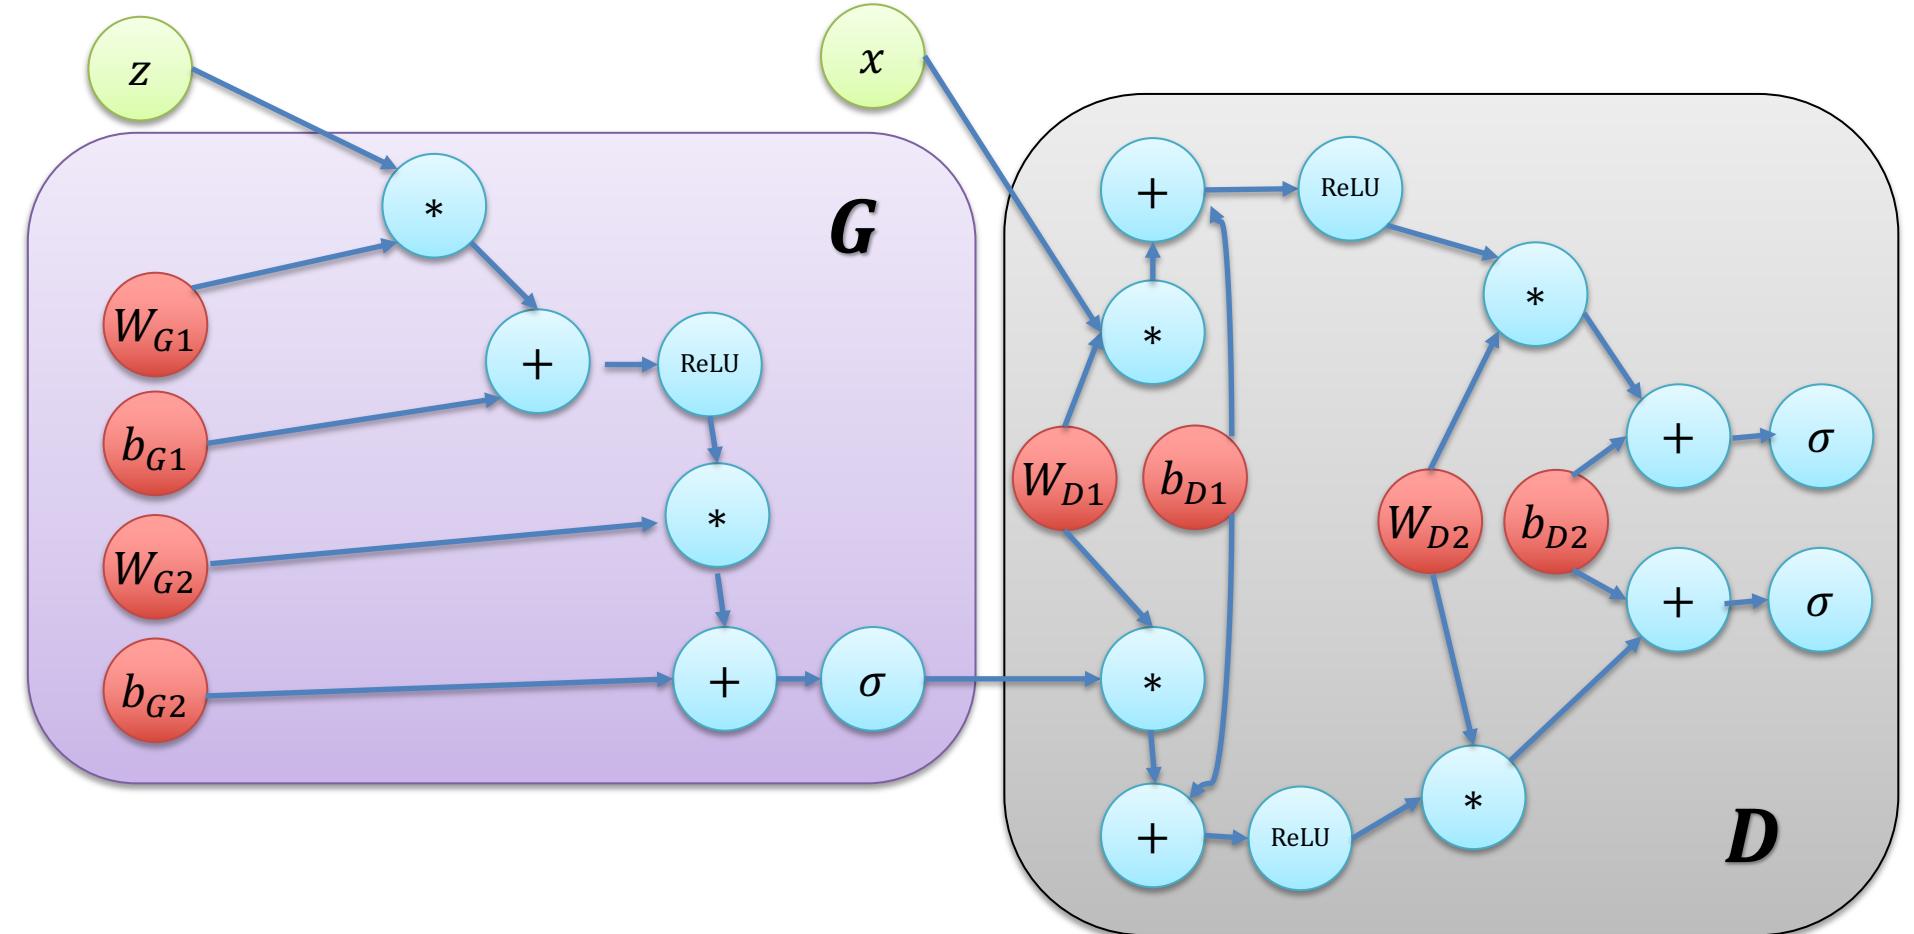
Architecture of D :

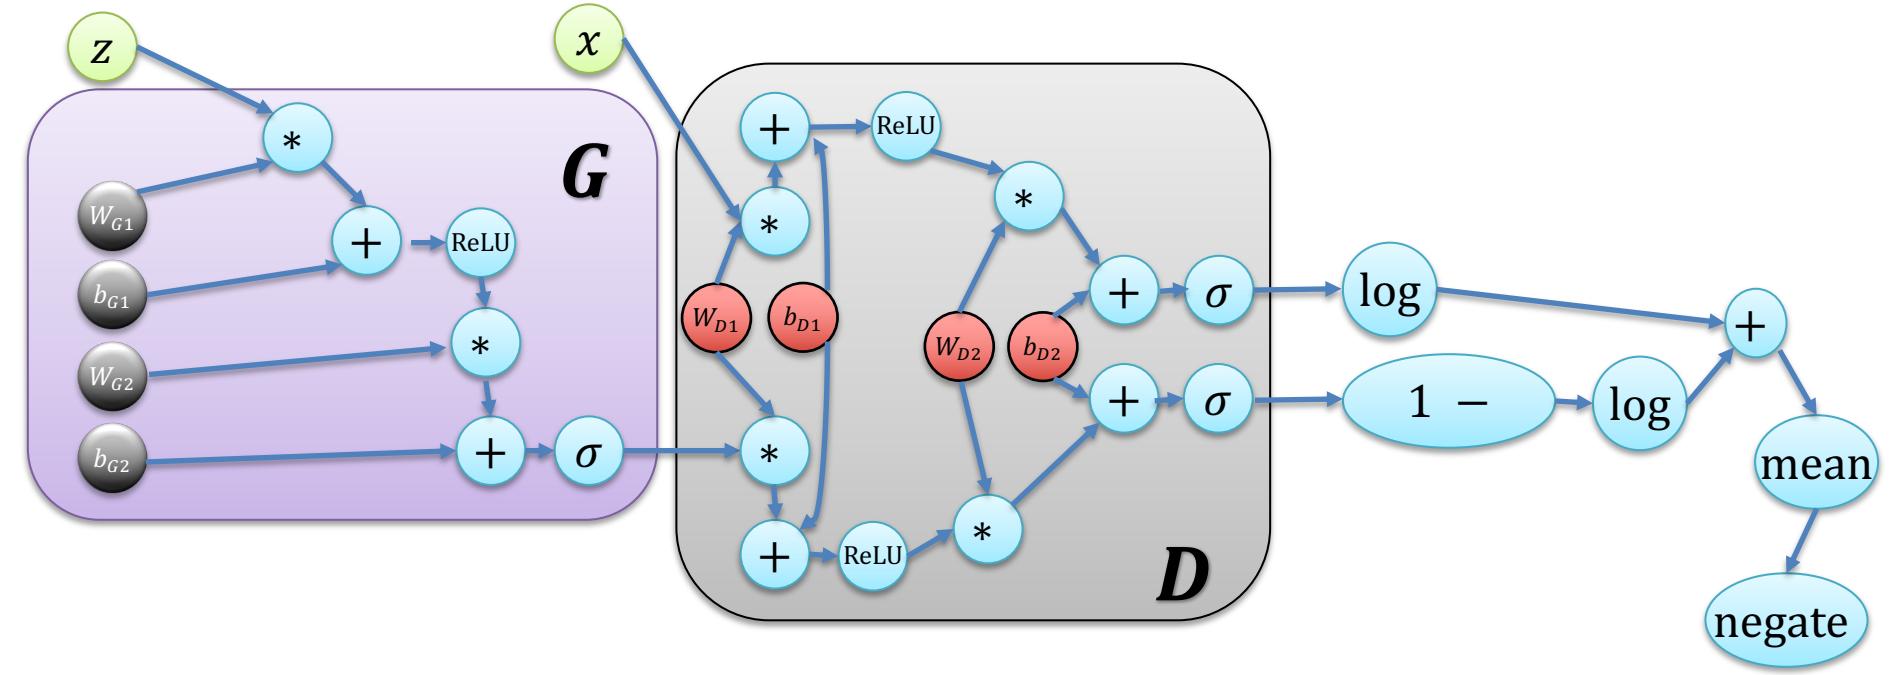
$784 \rightarrow 256 \rightarrow 1$

$W_{D1} = \text{tf.Variable}(\text{tf.truncated_normal}([784, 256], \text{stddev}=0.1))$

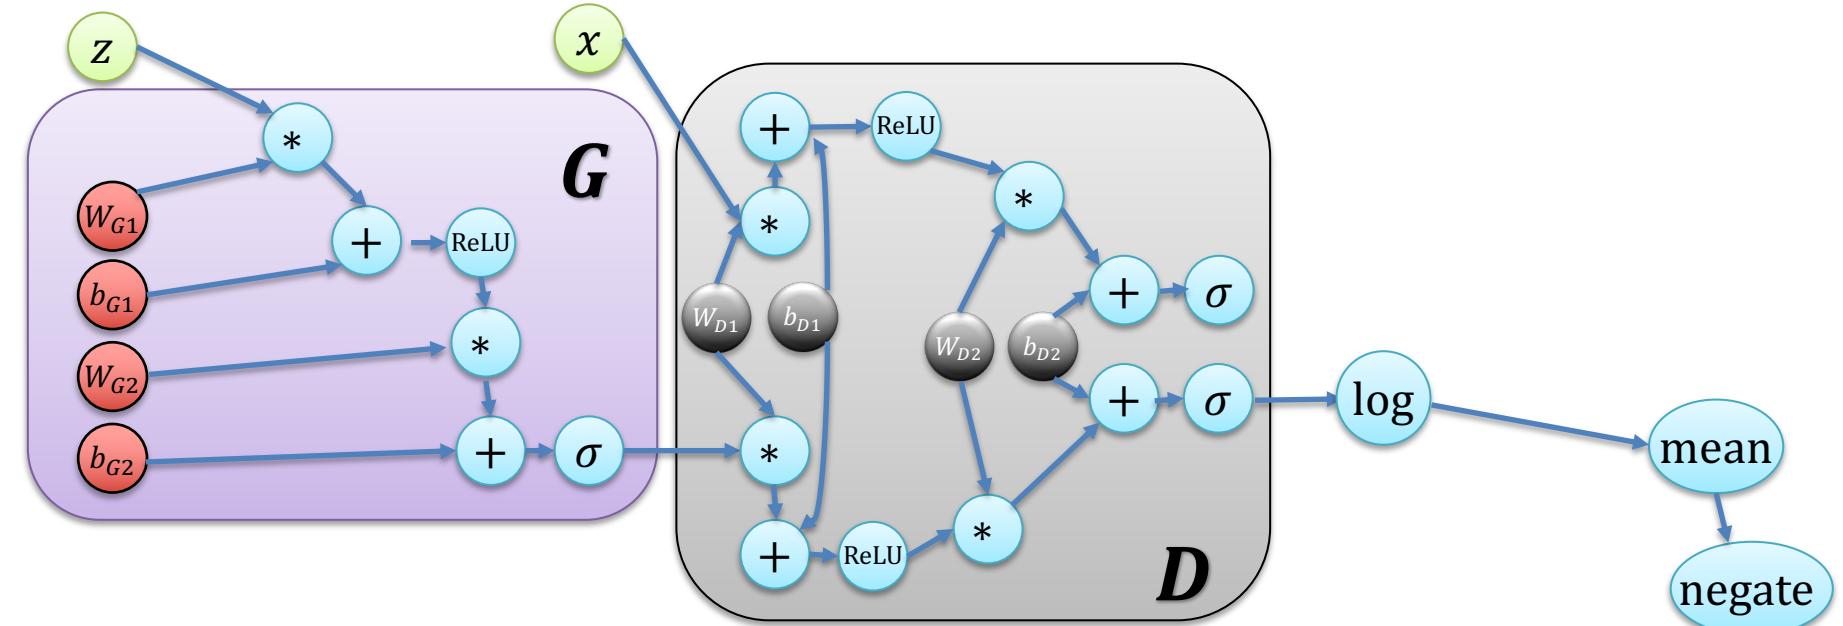


D

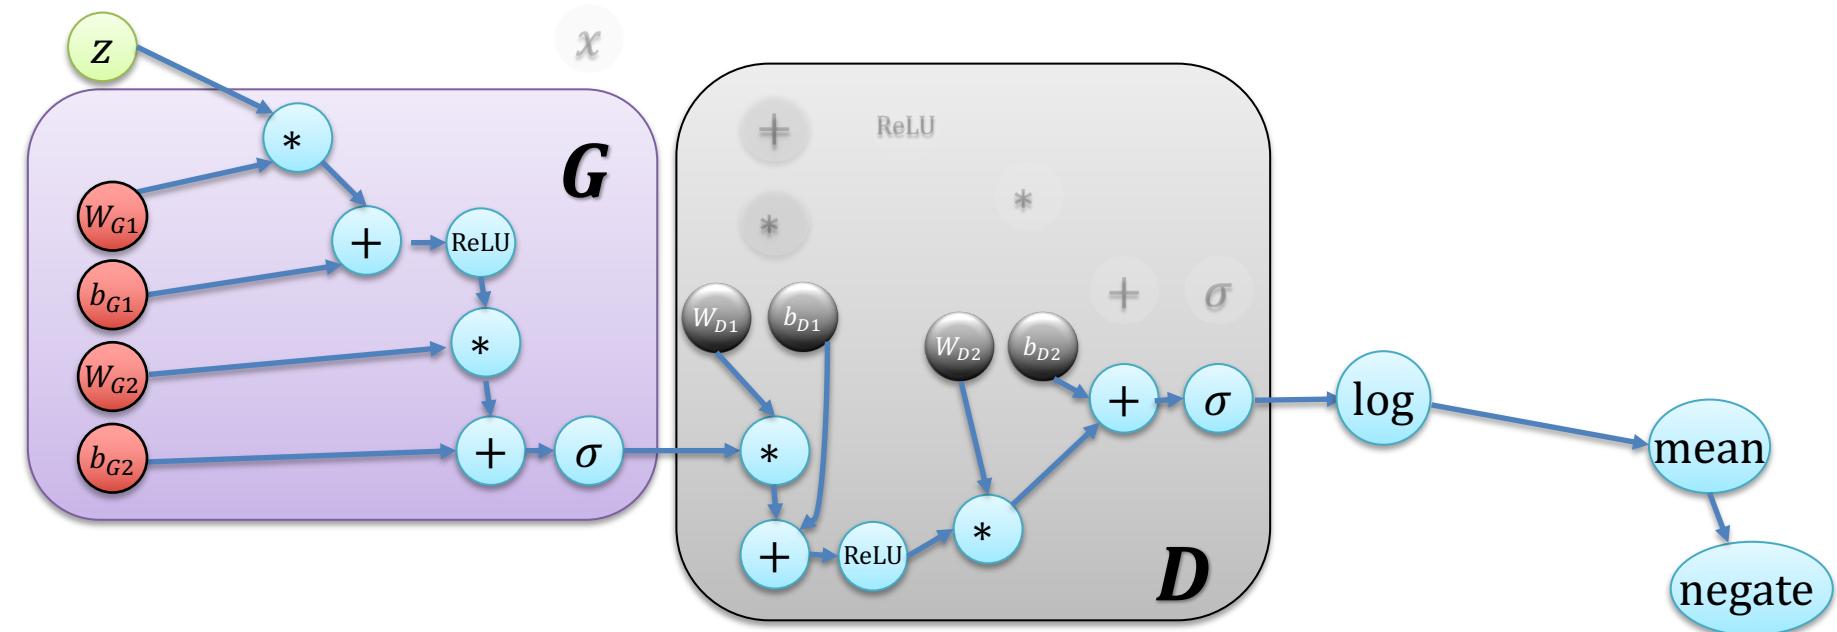




Loss when training **D**



Loss when training G



Loss when training G

